

# Deep Learning Homework 6 (Spring 2023)

This code is provided for Deep Learning class (601.482/682) Homework 6. For ease of implementation, we recommend working entire in Google Colaboratory.

@Copyright Cong Gao, the Johns Hopkins University, cgao11@jhu.edu. Modifications made by Hongtao Wu, Suzanna Sia, Hao Ding, Keith Harrigian, and Yiqing Shen.

## Imports

```
In [ ]: ## Mount Google Drive Data (If using Google Colaboratory)
try:
    from google.colab import drive
    drive.mount('/content/gdrive')
except:
    print("Mounting Failed.")
```

Mounted at /content/gdrive

```
In [ ]: ## Standard Library
import os
import json

## External Libraries
import numpy as np
import torch
import torch.nn as nn
from torchvision import transforms
from torch.autograd import Variable
import torch.nn.functional as functional
from torch.utils.data import Dataset, DataLoader
from skimage import io
import matplotlib.pyplot as plt
```

## Problem 1: Unsupervised Pre-training

## Training Hyperparameters

These are recommended hyperparameters - please feel free to use what works for you. Batch size can be changed if it does not match your memory, please state your batch\_step\_size in your report.

Dataset is available at: [https://livejohnshopkins-my.sharepoint.com/:u:/g/personal/yshen92\\_jh\\_edu/EcTxWAXsAhtDiv3vUxCTF8gBgAARCUvvKthb3s-pEExyMg](https://livejohnshopkins-my.sharepoint.com/:u:/g/personal/yshen92_jh_edu/EcTxWAXsAhtDiv3vUxCTF8gBgAARCUvvKthb3s-pEExyMg)

```
In [ ]: from google.colab import drive
drive.mount('/content/drive')
```

Mounted at /content/drive

```
In [ ]: ## Batch Size
train_batch_size = 10
validation_batch_size = 10

## Learning Rate
learning_rate = 0.001

# Epochs (Consider setting high and implementing early stopping)
num_epochs = 200
```

```
In [ ]: # from google.colab import drive
# drive.flush_and_unmount()
# drive.mount('/content/gdrive', force_remount=True)
```

Mounted at /content/gdrive

```
In [ ]: import os

# to check if the path already exist
if os.path.exists("/content/drive/MyDrive/HW6_data"):
    print("Directory exists. No need to unzip again.")
else:
    print("Directory not found. You might need to unzip again.")
```

Directory exists. No need to unzip again.

## Data Paths

```
In [ ]: # import zipfile

# with zipfile.ZipFile(data_dir, 'r') as zip_ref:
#     zip_ref.extractall("/content/gdrive/MyDrive/")

In [ ]: # General Data Directory ##TODO: Please fill in the appropriate directory
# data_dir = "/content/gdrive/MyDrive/hw6_data.zip"
data_dir = '/content/drive/MyDrive/HW6_data'

## Segmentation + Colorization Paths
segmentation_data_dir = f"{data_dir}/segmentation/"
colorization_data_dir = f"{data_dir}/colorization/"

# Mask JSON
mask_json = f"{data_dir}/mapping.json"

In [ ]: segmentation_data_dir

Out [ ]: '/content/drive/MyDrive/HW6_data/segmentation/'
```

## Data Loaders

We have provided you with some preprocessing code for the images but you should feel free to modify the class however you please to support your training schema. In the very least, you will have to modify the dataloader to support loading of the colorization dataset.

```
In [ ]: ## Image Transforms
img_transform = transforms.Compose([
    transforms.ToTensor(),
])

## Image Dataloader
class ImageDataset(Dataset):

    ....
    ImageDataset
```

```

#####

def __init__(self,
              input_dir,
              op,
              mask_json_path,
              transforms=None,
              is_colorization = False):
    #####
    ##TODO: Add support for colorization dataset

    Args:
        input_dir (str): Path to either colorization or segmentation directory
        op (str): One of "train", "val", or "test" signifying the desired split
        mask_json_path (str): Path to mapping.json file
        transforms (list or None): Image transformations to apply upon loading.
    #####

    self.transform = transforms
    self.op = op
    # Set the colorization flag
    self.is_colorization = is_colorization
    with open(mask_json_path, 'r') as f:
        self.mask = json.load(f)
    self.mask_num = len(self.mask) # There are 6 categories: grey, dark grey, and black
    self.mask_value = [int(value) for key, value in self.mask.items() if not key.startswith("_comment")]
    self.mask_value.sort()
    if self.is_colorization == False:
        if self.op == 'train':
            self.data_dir = os.path.join(input_dir, 'train')
        elif self.op == 'val':
            self.data_dir = os.path.join(input_dir, 'validation')
        elif self.op == 'test':
            self.data_dir = os.path.join(input_dir, 'test')
    else:
        if self.op == 'train':
            self.data_dir = os.path.join(input_dir, 'train_cor')
        elif self.op == 'val':
            self.data_dir = os.path.join(input_dir, 'validation_cor')

def __len__(self):
    #####

```

```

        return len(next(os.walk(self.data_dir))[1])

def __getitem__(self,
                idx):
    """

    """

    ## Load Image and Parse Properties
    if self.is_colorization == False:
        img_name = str(idx) + '_input.jpg'
        mask_name = str(idx) + '_mask.png'
    else:
        img_name = str(idx) + '_gray.jpg'
        mask_name = str(idx) + '_input.jpg'

    img = io.imread(os.path.join(self.data_dir, str(idx), img_name))
    mask = io.imread(os.path.join(self.data_dir, str(idx), mask_name))

    ## determine whether it is gray mask or RGB mask
    # there is no channel(c) for gray mask
    # c = 3 for RGB
    # h: height
    # w: width
    if len(mask.shape) == 2:
        h, w = mask.shape
    elif len(mask.shape) == 3:
        h, w, c = mask.shape
    ## Convert grey-scale label to one-hot encoding
    if self.is_colorization == False:
        new_mask = np.zeros((h, w, self.mask_num))
        for idx in range(self.mask_num):
            #if the mask has 3 dimension use this code
            new_mask[:, :, idx] = mask[:, :, 0] == self.mask_value[idx]
    else:
        new_mask = mask
        #if the mask has 1 dimension use the code below
        #new_mask[:, :, idx] = mask == self.mask_value[idx]
    ## Transform image and mask

```

```

    if self.transform:
        img, new_mask = self.img_transform(img, new_mask)
    # ## Use dictionary to output
    # sample = {'img': img, 'mask': mask}
    # return sample
    # 这里的mask已经转化为one-hot vector
    return img, new_mask

def img_transform(self,
                  img,
                  mask):
    """
    """
    """
    ## Apply Transformations to Image and Mask
    img = self.transform(img)
    mask = self.transform(mask)
    return img, mask

def img_transform_single(self, img):
    return self.transform(img)

```

## Model Architecture

Finish building the U-net architecture below.

```

In [ ]: ## Functions for adding the convolution layer
def add_conv_stage(dim_in,
                  dim_out,
                  kernel_size=3,
                  stride=1,
                  padding=1,
                  bias=True,
                  useBN=True):
    """
    """
    """
    # Use batch normalization
    if useBN:
        return nn.Sequential(

```

```

        nn.Conv2d(dim_in, dim_out, kernel_size=kernel_size, stride=stride, padding=padding, bias=bias),
        nn.BatchNorm2d(dim_out),
        nn.LeakyReLU(0.1),
        nn.Conv2d(dim_out, dim_out, kernel_size=kernel_size, stride=stride, padding=padding, bias=bias),
        nn.BatchNorm2d(dim_out),
        nn.LeakyReLU(0.1)
    )
    # No batch normalization
    else:
        return nn.Sequential(
            nn.Conv2d(dim_in, dim_out, kernel_size=kernel_size, stride=stride, padding=padding, bias=bias),
            nn.ReLU(),
            nn.Conv2d(dim_out, dim_out, kernel_size=kernel_size, stride=stride, padding=padding, bias=bias),
            nn.ReLU()
        )

## Upsampling
def upsample(ch_coarse,
             ch_fine):
    """
    """
    return nn.Sequential(
        nn.ConvTranspose2d(ch_coarse, ch_fine, 4, 2, 1, bias=False),
        nn.ReLU())

# U-Net
class UNET(nn.Module):
    """
    """
    def __init__(self, n_classes, input_channels_num = 3, useBN=True):
        """
        Args:
            n_classes (int): Number of classes
            useBN (bool): Turn Batch Norm on or off. (Hint: Using BatchNorm might help you achieve better
        """
        super(UNET, self).__init__()
        # Downgrade stages
        self.conv1 = add_conv_stage(input_channels_num, 32, useBN=useBN)

```

```

self.conv2 = add_conv_stage(32, 64, useBN=useBN)
self.conv3 = add_conv_stage(64, 128, useBN=useBN)
self.conv4 = add_conv_stage(128, 256, useBN=useBN)
# Upgrade stages
self.conv3m = add_conv_stage(256, 128, useBN=useBN)
self.conv2m = add_conv_stage(128, 64, useBN=useBN)
self.conv1m = add_conv_stage(64, 32, useBN=useBN)
# Maxpool
self.max_pool = nn.MaxPool2d(2)
# Upsample layers
self.upsample43 = upsample(256, 128)
self.upsample32 = upsample(128, 64)
self.upsample21 = upsample(64, 32)
# weight initialization
# You can have your own weight initialization. This is just an example.
for m in self.modules():
    if isinstance(m, nn.Conv2d) or isinstance(m, nn.ConvTranspose2d):
        if m.bias is not None:
            m.bias.data.zero_()
#TODO: Design your last layer & activations
self.final = nn.Conv2d(32, n_classes, 1)

if n_classes == 3:
    self.final_act = nn.Sigmoid()
else:
    self.final_act = nn.Softmax(dim = 1)

def forward(self, x):
    """
    Forward pass
    """
    conv1_out = self.conv1(x)
    conv2_out = self.conv2(self.max_pool(conv1_out))
    conv3_out = self.conv3(self.max_pool(conv2_out))
    conv4_out = self.conv4(self.max_pool(conv3_out))

    conv4m_out_ = torch.cat((self.upsample43(conv4_out), conv3_out), 1)
    conv3m_out = self.conv3m(conv4m_out_)

    conv3m_out_ = torch.cat((self.upsample32(conv3m_out), conv2_out), 1)
    conv2m_out = self.conv2m(conv3m_out_)

```



```

conv2m_out_ = torch.cat((self.upsample21(conv2m_out), conv1_out), 1)
conv1m_out = self.conv1m(conv2m_out_)

#TODO: Design your last layer & activations
final_out = self.final(conv1m_out)

return self.final_act(final_out)

```

```

In [ ]: device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
print(f"Current device: {device}")

```

Current device: cuda

```

In [ ]: img = io.imread('/content/drive/MyDrive/HW6_data/colorization/validation_cor/1/1_gray.jpg')
print(img.shape)

```

(256, 320)

## DICE Score and DICE Loss

Finish implementing the DICE score function below and then write a Dice Loss function that you can use to update your model weights.

```

In [ ]: ##TODO: Finish implementing the multi-class DICE score function
def dice_score_image(prediction, target, n_classes):
    """
    computer the mean dice score for a single image

    Reminders: A false positive is a result that indicates a given condition exists, when it does not
               A false negative is a test result that indicates that a condition does not hold, while in fact it does
    Args:
        prediction (tensor): predicted labels of the image
        target (tensor): ground truth of the image
        n_classes (int): number of classes

    Returns:
        m_dice (float): Mean dice score over classes
    """
    ## Should test image one by one
    assert prediction.shape[0] == 1 #This line can not be deleted

```

```

## TODO: Compute Dice Score for Each Class. Compute Mean Dice Score over Classes.
dice_scores = np.zeros(n_classes)
target = torch.argmax(target, dim = 1)
for cl in range(n_classes):
    pred_cl = (prediction == cl).float()
    true_cl = (target == cl).float()
    TP = torch.sum(pred_cl * true_cl)
    FP = torch.sum(pred_cl*(1 - true_cl))
    FN = torch.sum((1 - pred_cl) * true_cl)
    epsilon = 1e-7
    #When there is no ground truth of the class in this image
    #Give 1 dice score if False Positive pixel number is 0,
    #give 0 dice score if False Positive pixel number is not 0 (> 0).
    # if FP == 0:
    #     dice_scores.append(1.0)
    # else:
    #     dice_scores.append( (2 * TP + epsilon) / (2 * TP + FP + FN + epsilon) )

    if TP == 0 or FN == 0:
        dice_scores[cl] = 1.0 if FP == 0 else 0.0
    else:
        dice_scores[cl] = (2 * TP + epsilon) / (2 * TP + FP + FN + epsilon)

return dice_scores.mean()

```

**def** dice\_score\_dataset(model, dataloader, num\_classes, use\_gpu=False):

"""

Compute the mean dice score on a set of data.

Note that multiclass dice score can be defined as the mean over classes of binary dice score. Dice score is computed per image. Mean dice score over the dataset is the dice score averaged across all images.

Reminders: A false positive is a result that indicates a given condition exists, when it does not  
A false negative is a test result that indicates that a condition does not hold, while in fact it does.

Args:

- model (UNET class): Your trained model
- dataloader (DataLoader): Dataset for evaluation
- num\_classes (int): Number of classes

```

Returns:
    m_dice (float): Mean dice score over the input dataset
    """
    ## Number of Batches and Cache over Dataset
    n_batches = len(data_loader)
    scores = np.zeros(n_batches)
    ## Evaluate
    model.eval()
    idx = 0
    with torch.no_grad():
        for data in data_loader:
            ## Format Data
            img, target = data
            if use_gpu:
                img = img.cuda()
                target = target.cuda()
            ## Make Predictions
            out = model(img)
            n_classes = out.shape[1]
            prediction = torch.argmax(out, dim = 1)
            scores[idx] = dice_score_image(prediction, target, n_classes)
            idx += 1
    ## Average Dice Score Over Images
    m_dice = scores.mean()
    return m_dice

## TODO: Implement DICE loss,
# It should conform to to how we computer the dice score.
class DICELoss(nn.Module):
    def __init__(self, n_classes):
        super(DICELoss, self).__init__()
        self.n_classes = n_classes
    def forward(self, prediction, target):
        #probs = torch.nn.functional.softmax(prediction, dim=1)
        probs = prediction
        dice_loss = 0.0
        for cl in range(self.n_classes):
            #pred_cl = (prediciton_classes == cl).float()
            # true_cl = (target == cl).float()
            pred_cl = probs[:, cl, :, :]

```

```

    true_cl = target[:, cl, :, :]
    error = 1e-7

    TP = torch.sum(pred_cl * true_cl)
    FP = torch.sum(pred_cl * (1 - true_cl))
    FN = torch.sum((1 - pred_cl) * true_cl)

    # Add dice loss for the class to the total dice loss
    dice_score = (2 * TP + error) / (2 * TP + FP + FN + error)
    dice_loss += 1 - dice_score
    return dice_loss / self.n_classes

```

```

In [ ]: import random

def set_seed(seed_value):
    """Set seed for reproducibility."""
    random.seed(seed_value)
    np.random.seed(seed_value)
    torch.manual_seed(seed_value)
    torch.cuda.manual_seed(seed_value)
    torch.cuda.manual_seed_all(seed_value)
    torch.backends.cudnn.deterministic = True
    torch.backends.cudnn.benchmark = False
seed_value = 12345
set_seed(seed_value)

```

## Training Procedure (Segmentation)

```

In [ ]: import copy

```

```

In [ ]: ## Initialize your unet
n_classes = len(json.load(open(mask_json)))
model = UNET(n_classes)
device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
model = model.to(device)

## Initialize Dataloaders
train_dataset=ImageDataset(input_dir=segmentation_data_dir, op="train", mask_json_path=mask_json, transform=
validation_dataset=ImageDataset(input_dir=segmentation_data_dir, op="val", mask_json_path=mask_json, transform=
test_dataset=ImageDataset(input_dir=segmentation_data_dir, op="test", mask_json_path=mask_json, transforms=

```

```

train_dataloader = DataLoader(train_dataset, batch_size=train_batch_size, shuffle=True)
validation_dataloader = DataLoader(validation_dataset, batch_size=validation_batch_size, shuffle=False)
test_dataloader = DataLoader(test_dataset, batch_size=1, shuffle=False)
## Initialize Optimizer and Learning Rate Scheduler
optimizer = torch.optim.Adam(model.parameters(), lr=learning_rate)
scheduler = torch.optim.lr_scheduler.StepLR(optimizer, step_size=10, gamma=0.1)

# implement early stopping
patience = 10
best_val_loss = float('inf')
counter_early_stop = 0
early_stop = False
training_loss_values = []
validation_loss_values = []
print("Start Training...")
for epoch in range(num_epochs):
    ##### Training #####
    print("\nEPOCH " + str(epoch+1) + " of " + str(num_epochs) + "\n")
    # TODO: Design your own training section
    model.train()
    train_loss = 0.0
    for images, masks in train_dataloader:
        images = images.to(device)
        masks = masks.to(device)

        optimizer.zero_grad()
        outputs = model(images)
        # print(outputs.requires_grad)
        loss = DICELoss(n_classes)(outputs, masks)
        loss.backward()
        optimizer.step()
        train_loss += loss.item()

    train_loss = train_loss / len(train_dataloader)
    training_loss_values.append(train_loss)

    ##### Validation #####
    # TODO: Design your own validation section
    model.eval()
    val_loss = 0.0
    with torch.no_grad():
        for images, masks in validation_dataloader:

```

```

        images = images.to(device)
        masks = masks.to(device)
        outputs = model(images)
        # probs = torch.softmax(outputs, dim = 1)
        loss = DICELoss(n_classes).forward(outputs, masks)
        val_loss += loss.item()
    val_loss = val_loss / len(validation_dataloader)
    validation_loss_values.append(val_loss)
    print(f"Epoch {epoch+1}, Training loss: {train_loss:.4f}, Validation Loss: {val_loss:.4f}")
    if val_loss < best_val_loss:
        best_val_loss = val_loss
        best_model = copy.deepcopy(model.state_dict()) # Save the best model
        counter_early_stop = 0 # Reset counter
        print("Validation loss decreased, saving model...")
    else:
        counter_early_stop += 1
        print(f"Validation loss did not decrease, counter: {counter_early_stop}/{patience}")

    if counter_early_stop >= patience or epoch == num_epochs - 1:
        early_stop = True
        break

    scheduler.step()

plt.figure(figsize=(10, 5))
plt.plot(training_loss_values, label='Training Loss')
plt.plot(validation_loss_values, label='Validation Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.title('Loss vs. Epochs')
plt.legend()
plt.show()

```

Start Training...

EPOCH 1 of 200

Epoch 1, Training loss: 0.8444, Validation Loss: 0.8090  
Validation loss decreased, saving model...

EPOCH 2 of 200

Epoch 2, Training loss: 0.7666, Validation Loss: 0.8068  
Validation loss decreased, saving model...

EPOCH 3 of 200

Epoch 3, Training loss: 0.7158, Validation Loss: 0.7713  
Validation loss decreased, saving model...

EPOCH 4 of 200

Epoch 4, Training loss: 0.6663, Validation Loss: 0.7274  
Validation loss decreased, saving model...

EPOCH 5 of 200

Epoch 5, Training loss: 0.6166, Validation Loss: 0.8029  
Validation loss did not decrease, counter: 1/10

EPOCH 6 of 200

Epoch 6, Training loss: 0.6155, Validation Loss: 0.6704  
Validation loss decreased, saving model...

EPOCH 7 of 200

Epoch 7, Training loss: 0.5525, Validation Loss: 0.6612  
Validation loss decreased, saving model...

EPOCH 8 of 200

Epoch 8, Training loss: 0.5449, Validation Loss: 0.6834  
Validation loss did not decrease, counter: 1/10

EPOCH 9 of 200

Epoch 9, Training loss: 0.5215, Validation Loss: 0.6658  
Validation loss did not decrease, counter: 2/10

EPOCH 10 of 200

Epoch 10, Training loss: 0.4923, Validation Loss: 0.6545  
Validation loss decreased, saving model...

EPOCH 11 of 200

Epoch 11, Training loss: 0.4636, Validation Loss: 0.6272  
Validation loss decreased, saving model...

EPOCH 12 of 200

Epoch 12, Training loss: 0.4316, Validation Loss: 0.6183  
Validation loss decreased, saving model...

EPOCH 13 of 200

Epoch 13, Training loss: 0.4274, Validation Loss: 0.6157  
Validation loss decreased, saving model...

EPOCH 14 of 200

Epoch 14, Training loss: 0.4171, Validation Loss: 0.6175  
Validation loss did not decrease, counter: 1/10

EPOCH 15 of 200

Epoch 15, Training loss: 0.3945, Validation Loss: 0.6139  
Validation loss decreased, saving model...

EPOCH 16 of 200

Epoch 16, Training loss: 0.4149, Validation Loss: 0.6140  
Validation loss did not decrease, counter: 1/10

EPOCH 17 of 200



Epoch 17, Training loss: 0.4108, Validation Loss: 0.6079  
Validation loss decreased, saving model...

EPOCH 18 of 200

Epoch 18, Training loss: 0.4037, Validation Loss: 0.6090  
Validation loss did not decrease, counter: 1/10

EPOCH 19 of 200

Epoch 19, Training loss: 0.3898, Validation Loss: 0.6080  
Validation loss did not decrease, counter: 2/10

EPOCH 20 of 200

Epoch 20, Training loss: 0.3694, Validation Loss: 0.6095  
Validation loss did not decrease, counter: 3/10

EPOCH 21 of 200

Epoch 21, Training loss: 0.3827, Validation Loss: 0.6034  
Validation loss decreased, saving model...

EPOCH 22 of 200

Epoch 22, Training loss: 0.3801, Validation Loss: 0.6037  
Validation loss did not decrease, counter: 1/10

EPOCH 23 of 200

Epoch 23, Training loss: 0.3939, Validation Loss: 0.6027  
Validation loss decreased, saving model...

EPOCH 24 of 200

Epoch 24, Training loss: 0.3894, Validation Loss: 0.6024  
Validation loss decreased, saving model...

EPOCH 25 of 200

Epoch 25, Training loss: 0.3848, Validation Loss: 0.6024  
Validation loss decreased, saving model...

EPOCH 26 of 200

Epoch 26, Training loss: 0.3967, Validation Loss: 0.6013  
Validation loss decreased, saving model...

EPOCH 27 of 200

Epoch 27, Training loss: 0.3722, Validation Loss: 0.6022  
Validation loss did not decrease, counter: 1/10

EPOCH 28 of 200

Epoch 28, Training loss: 0.3772, Validation Loss: 0.6018  
Validation loss did not decrease, counter: 2/10

EPOCH 29 of 200

Epoch 29, Training loss: 0.3900, Validation Loss: 0.6022  
Validation loss did not decrease, counter: 3/10

EPOCH 30 of 200

Epoch 30, Training loss: 0.3755, Validation Loss: 0.6007  
Validation loss decreased, saving model...

EPOCH 31 of 200

Epoch 31, Training loss: 0.3698, Validation Loss: 0.6012  
Validation loss did not decrease, counter: 1/10

EPOCH 32 of 200

Epoch 32, Training loss: 0.3628, Validation Loss: 0.6009  
Validation loss did not decrease, counter: 2/10

EPOCH 33 of 200

Epoch 33, Training loss: 0.3621, Validation Loss: 0.6019  
Validation loss did not decrease, counter: 3/10

EPOCH 34 of 200

Epoch 34, Training loss: 0.3815, Validation Loss: 0.6007  
Validation loss did not decrease, counter: 4/10

EPOCH 35 of 200

Epoch 35, Training loss: 0.3744, Validation Loss: 0.6013  
Validation loss did not decrease, counter: 5/10

EPOCH 36 of 200

Epoch 36, Training loss: 0.3759, Validation Loss: 0.6009  
Validation loss did not decrease, counter: 6/10

EPOCH 37 of 200

Epoch 37, Training loss: 0.3741, Validation Loss: 0.6021  
Validation loss did not decrease, counter: 7/10

EPOCH 38 of 200


Epoch 38, Training loss: 0.3878, Validation Loss: 0.6011  
Validation loss did not decrease, counter: 8/10

EPOCH 39 of 200

Epoch 39, Training loss: 0.3802, Validation Loss: 0.6007  
Validation loss did not decrease, counter: 9/10

EPOCH 40 of 200

Epoch 40, Training loss: 0.3659, Validation Loss: 0.6007  
Validation loss did not decrease, counter: 10/10



```
In [ ]: test_dice_score = dice_score_dataset(model, test_dataloader, n_classes, use_gpu=device.type == 'cuda')
print(f"Epoch {epoch+1}, DICE score on the test set: {test_dice_score:.4f}")
```

Epoch 40, DICE score on the test set: 0.5527

```
In [ ]: from torchvision.transforms import functional as F
```

```
class RandomHorizontalFlip(object):
```

```
    def __init__(self, prob):
        self.prob = prob
```

```
    def __call__(self, img):
        if torch.rand(1).item() < self.prob:
```

```

        return torch.flip(img, [2])
    return img

class RandomRotation(object):

    def __init__(self, degrees):
        self.degrees = degrees

    def __call__(self, img):
        angle = torch.FloatTensor(1).uniform_(-self.degrees, self.degrees).item()
        return F.rotate(img, angle)

##prob = 0.25, degrees = 25, dice score = 0.5623

train_img_transform= transforms.Compose([
    transforms.ToTensor(),
    RandomHorizontalFlip (prob=0.25),
    RandomRotation (degrees=15),
])

```

```

In [ ]: ## Initialize your unet
n_classes = len(json.load(open(mask_json)))
model = UNET(n_classes)
device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
model = model.to(device)

## Initialize Dataloaders
train_dataset=ImageDataset(input_dir=segmentation_data_dir, op="train", mask_json_path=mask_json, transform=
validation_dataset=ImageDataset(input_dir=segmentation_data_dir, op="val", mask_json_path=mask_json, transform=
test_dataset=ImageDataset(input_dir=segmentation_data_dir, op="test", mask_json_path=mask_json, transforms=
train_dataloader = DataLoader(train_dataset, batch_size=train_batch_size, shuffle=True)
validation_dataloader = DataLoader(validation_dataset, batch_size=validation_batch_size, shuffle=False)
test_dataloader = DataLoader(test_dataset, batch_size=1, shuffle=False)
## Initialize Optimizer and Learning Rate Scheduler
optimizer = torch.optim.Adam(model.parameters(),lr=learning_rate)
scheduler = torch.optim.lr_scheduler.StepLR(optimizer, step_size=10, gamma=0.1)

# implement early stopping

```

```

patience = 10 # How many epochs to wait after last time validation loss improved.
best_val_loss = float('inf')
counter_early_stop = 0
early_stop = False
training_loss_values = []
validation_loss_values = []
print("Start Training...")
for epoch in range(num_epochs):
    ##### Training #####
    print("\nEPOCH " + str(epoch+1) + " of " + str(num_epochs) + "\n")
    # TODO: Design your own training section
    model.train()
    train_loss = 0.0

    for images, masks in train_dataloader:

        # if torch.rand(1).item() < 0.5:
        #     images = torch.flip(images, [2])
        #     masks = torch.flip(masks, [2])
        # if torch.rand(1).item() < 0.5:
        #     images = torch.flip(images, [3])
        #     masks = torch.flip(masks, [3])

        images = images.to(device)
        masks = masks.to(device)
        optimizer.zero_grad()
        outputs = model(images)
        # print(outputs.requires_grad)
        loss = DICELoss(n_classes)(outputs, masks)
        loss.backward()
        optimizer.step()
        train_loss += loss.item()

    train_loss = train_loss / len(train_dataloader)
    training_loss_values.append(train_loss)

    ##### Validation #####
    # TODO: Design your own validation section
    model.eval()
    val_loss = 0.0
    with torch.no_grad():

```

```

    for images, masks in validation_dataloader:
        images = images.to(device)
        masks = masks.to(device)
        outputs = model(images)
        # probs = torch.softmax(outputs, dim = 1)
        loss = DICELoss(n_classes).forward(outputs, masks)
        val_loss += loss.item()
    val_loss = val_loss / len(validation_dataloader)
    validation_loss_values.append(val_loss)
    print(f"Epoch {epoch+1}, Training loss: {train_loss:.4f}, Validation Loss: {val_loss:.4f}")
    if val_loss < best_val_loss:
        best_val_loss = val_loss
        best_model = copy.deepcopy(model.state_dict()) # Save the best model
        counter_early_stop = 0 # Reset counter
        print("Validation loss decreased, saving model...")
    else:
        counter_early_stop += 1
        print(f"Validation loss did not decrease, counter: {counter_early_stop}/{patience}")

    if counter_early_stop >= patience or epoch == num_epochs - 1:
        early_stop = True
        break

    scheduler.step()

plt.figure(figsize=(10, 5))
plt.plot(training_loss_values, label='Training Loss')
plt.plot(validation_loss_values, label='Validation Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.title('Loss vs. Epochs')
plt.legend()
plt.show()

```

Start Training...

EPOCH 1 of 200

Epoch 1, Training loss: 0.8672, Validation Loss: 0.8139  
Validation loss decreased, saving model...

EPOCH 2 of 200

Epoch 2, Training loss: 0.8249, Validation Loss: 0.7752  
Validation loss decreased, saving model...

EPOCH 3 of 200

Epoch 3, Training loss: 0.7948, Validation Loss: 0.7762  
Validation loss did not decrease, counter: 1/10

EPOCH 4 of 200

Epoch 4, Training loss: 0.7779, Validation Loss: 0.7761  
Validation loss did not decrease, counter: 2/10

EPOCH 5 of 200

Epoch 5, Training loss: 0.7647, Validation Loss: 0.7601  
Validation loss decreased, saving model...

EPOCH 6 of 200

Epoch 6, Training loss: 0.7436, Validation Loss: 0.7867  
Validation loss did not decrease, counter: 1/10

EPOCH 7 of 200

Epoch 7, Training loss: 0.7326, Validation Loss: 0.7897  
Validation loss did not decrease, counter: 2/10

EPOCH 8 of 200

Epoch 8, Training loss: 0.7401, Validation Loss: 0.7948  
Validation loss did not decrease, counter: 3/10



EPOCH 9 of 200

Epoch 9, Training loss: 0.7265, Validation Loss: 0.7869  
Validation loss did not decrease, counter: 4/10

EPOCH 10 of 200

Epoch 10, Training loss: 0.7263, Validation Loss: 0.7745  
Validation loss did not decrease, counter: 5/10

EPOCH 11 of 200

Epoch 11, Training loss: 0.7127, Validation Loss: 0.6979  
Validation loss decreased, saving model...

EPOCH 12 of 200

Epoch 12, Training loss: 0.7119, Validation Loss: 0.7036  
Validation loss did not decrease, counter: 1/10

EPOCH 13 of 200

Epoch 13, Training loss: 0.7141, Validation Loss: 0.6890  
Validation loss decreased, saving model...

EPOCH 14 of 200

Epoch 14, Training loss: 0.7035, Validation Loss: 0.6856  
Validation loss decreased, saving model...

EPOCH 15 of 200

Epoch 15, Training loss: 0.7077, Validation Loss: 0.6920  
Validation loss did not decrease, counter: 1/10

EPOCH 16 of 200

Epoch 16, Training loss: 0.7143, Validation Loss: 0.6833  
Validation loss decreased, saving model...

EPOCH 17 of 200

Epoch 17, Training loss: 0.6956, Validation Loss: 0.6788  
Validation loss decreased, saving model...

EPOCH 18 of 200

Epoch 18, Training loss: 0.7073, Validation Loss: 0.6734  
Validation loss decreased, saving model...

EPOCH 19 of 200

Epoch 19, Training loss: 0.7028, Validation Loss: 0.6797  
Validation loss did not decrease, counter: 1/10

EPOCH 20 of 200

Epoch 20, Training loss: 0.6883, Validation Loss: 0.6733  
Validation loss decreased, saving model...

EPOCH 21 of 200

Epoch 21, Training loss: 0.7126, Validation Loss: 0.6712  
Validation loss decreased, saving model...

EPOCH 22 of 200

Epoch 22, Training loss: 0.7006, Validation Loss: 0.6714  
Validation loss did not decrease, counter: 1/10

EPOCH 23 of 200

Epoch 23, Training loss: 0.7025, Validation Loss: 0.6731  
Validation loss did not decrease, counter: 2/10

EPOCH 24 of 200

Epoch 24, Training loss: 0.6826, Validation Loss: 0.6716  
Validation loss did not decrease, counter: 3/10

EPOCH 25 of 200

Epoch 25, Training loss: 0.7120, Validation Loss: 0.6702  
Validation loss decreased, saving model...

EPOCH 26 of 200

Epoch 26, Training loss: 0.6892, Validation Loss: 0.6707  
Validation loss did not decrease, counter: 1/10

EPOCH 27 of 200

Epoch 27, Training loss: 0.7057, Validation Loss: 0.6716  
Validation loss did not decrease, counter: 2/10

EPOCH 28 of 200

Epoch 28, Training loss: 0.6857, Validation Loss: 0.6706  
Validation loss did not decrease, counter: 3/10

EPOCH 29 of 200

Epoch 29, Training loss: 0.7031, Validation Loss: 0.6707  
Validation loss did not decrease, counter: 4/10

EPOCH 30 of 200

Epoch 30, Training loss: 0.6933, Validation Loss: 0.6722  
Validation loss did not decrease, counter: 5/10

EPOCH 31 of 200

Epoch 31, Training loss: 0.6906, Validation Loss: 0.6717  
Validation loss did not decrease, counter: 6/10

EPOCH 32 of 200

Epoch 32, Training loss: 0.6931, Validation Loss: 0.6731  
Validation loss did not decrease, counter: 7/10

EPOCH 33 of 200


Epoch 33, Training loss: 0.6819, Validation Loss: 0.6715  
Validation loss did not decrease, counter: 8/10

EPOCH 34 of 200

Epoch 34, Training loss: 0.6789, Validation Loss: 0.6708  
Validation loss did not decrease, counter: 9/10

EPOCH 35 of 200

Epoch 35, Training loss: 0.6830, Validation Loss: 0.6717  
Validation loss did not decrease, counter: 10/10



```
In [ ]: test_dice_score = dice_score_dataset(model, test_dataloader, n_classes, use_gpu=device.type == 'cuda')
        print(f"Epoch {epoch+1}, DICE score on the test set: {test_dice_score:.4f}")
```

Epoch 35, DICE score on the test set: 0.4572

# Training Procedure: Colorization Pre-training

Complete the rest of this problem in the cells below.

```
In [ ]: ## Batch Size
train_batch_size = 30
validation_batch_size = 30

## Learning Rate
learning_rate = 0.001

# Epochs (Consider setting high and implementing early stopping)
num_epochs = 200
```

```
In [ ]: ## Image Transforms
img_transform = transforms.Compose([
    transforms.ToTensor(),
])

## Image Dataloader
class ImageDataset(Dataset):

    """
    ImageDataset
    """

    def __init__(self,
                 input_dir,
                 op,
                 mask_json_path,
                 transforms=None,
                 is_colorization = False):

        """
        ##TODO: Add support for colorization dataset

        Args:
            input_dir (str): Path to either colorization or segmentation directory
            op (str): One of "train", "val", or "test" signifying the desired split
            mask_json_path (str): Path to mapping.json file
```

```

        transforms (list or None): Image transformations to apply upon loading.
    """
    self.transform = transforms
    self.op = op
    # Set the colorization flag
    self.is_colorization = is_colorization
    with open(mask_json_path, 'r') as f:
        self.mask = json.load(f)
    self.mask_num = len(self.mask) # There are 6 categories: grey, dark grey, and black
    self.mask_value = [int(value) for key, value in self.mask.items() if not key.startswith("_comment")]
    self.mask_value.sort()
    if self.is_colorization == False:
        if self.op == 'train':
            self.data_dir = os.path.join(input_dir, 'train')
        elif self.op == 'val':
            self.data_dir = os.path.join(input_dir, 'validation')
        elif self.op == 'test':
            self.data_dir = os.path.join(input_dir, 'test')
    else:
        if self.op == 'train':
            self.data_dir = os.path.join(input_dir, 'train_cor')
        elif self.op == 'val':
            self.data_dir = os.path.join(input_dir, 'validation_cor')

    def __len__(self):
        if self.op == 'train':
            return 1584
        elif self.op == 'val':
            return 50
        else:
            raise ValueError("Invalid split option: {}".format(self.op))

    def __getitem__(self,
                    idx):
        """
        """

        ## Load Image and Parse Properties
        if self.is_colorization == False:
            img_name = str(idx) + '_input.jpg'
            mask_name = str(idx) + '_mask.png'

```

```

else:
    img_name = str(idx) + '_gray.jpg'
    mask_name = str(idx) + '_input.jpg'

img = io.imread(os.path.join(self.data_dir, str(idx), img_name))
mask = io.imread(os.path.join(self.data_dir, str(idx), mask_name))

## determine whether it is gray mask or RGB mask
# there is no channel(c) for gray mask
# c = 3 for RGB
# h: height
# w: width
if len(mask.shape) == 2:
    h, w = mask.shape
elif len(mask.shape) == 3:
    h, w, c = mask.shape
## Convert grey-scale label to one-hot encoding
if self.is_colorization == False:
    new_mask = np.zeros((h, w, self.mask_num))
    for idx in range(self.mask_num):
        #if the mask has 3 dimension use this code
        new_mask[:, :, idx] = mask[:, :, 0] == self.mask_value[idx]
else:
    new_mask = mask
    #if the mask has 1 dimension use the code below
    #new_mask[:, :, idx] = mask == self.mask_value[idx]
## Transform image and mask
if self.transform:
    img, new_mask = self.img_transform(img, new_mask)
# ## Use dictionary to output
# sample = {'img': img, 'mask': mask}
# return sample
# 这里的mask已经转化为one-hot vector
return img, new_mask

def img_transform(self,
                  img,
                  mask):
    """

```

```

        #####
        ## Apply Transformations to Image and Mask
        img = self.transform(img)
        mask = self.transform(mask)
        return img, mask

    def img_transform_single(self, img):
        return self.transform(img)

```

```

In [ ]: from torch.optim.lr_scheduler import StepLR
        from torch.nn import MSELoss

        # for colorization task, we expect the output channel(n_classes) = 3
        model = UNET(input_channels_num= 1,n_classes = 3)
        device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
        model = model.to(device)

        ## Initialize Dataloaders
        train_dataset=ImageDataset(input_dir=colorization_data_dir, op="train", mask_json_path=mask_json, transform=transform)
        validation_dataset=ImageDataset(input_dir=colorization_data_dir, op="val", mask_json_path=mask_json, transform=transform)
        train_dataloader = DataLoader(train_dataset, batch_size=train_batch_size, shuffle=True)
        validation_dataloader = DataLoader(validation_dataset, batch_size=validation_batch_size, shuffle=False)

        ## Initialize Optimizer and Learning Rate Scheduler
        optimizer = torch.optim.Adam(model.parameters(),lr=learning_rate)
        scheduler = torch.optim.lr_scheduler.StepLR(optimizer, step_size=10, gamma=0.1)
        criterion = nn.MSELoss()

        patience = 5 # How many epochs to wait after last time validation loss improved.
        best_val_loss = float('inf')
        counter_early_stop = 0
        early_stop = False
        training_loss_values = []
        validation_loss_values = []
        print("Start Training...")
        for epoch in range(num_epochs):
            ##### Training #####
            print("\nEPOCH " +str(epoch+1)+" of "+str(num_epochs)+"\n")
            # TODO: Design your own training section
            model.train()
            train_loss = 0.0
            for images, mask in train_dataloader: # No masks needed for colorization
                images = images.to(device)

```



```

mask = mask.to(device)

outputs = model(images)
loss = criterion(outputs, mask)

optimizer.zero_grad()
loss.backward()
optimizer.step()
train_loss += loss.item()

train_loss = train_loss / len(train_dataloader)
training_loss_values.append(train_loss)

##### Validation #####
# TODO: Design your own validation section
model.eval()
val_loss = 0.0
with torch.no_grad():
    for images, masks in validation_dataloader:
        images = images.to(device)
        masks = masks.to(device)

        outputs = model(images)
        loss = criterion(outputs, masks)

        val_loss += loss.item()

val_loss = val_loss / len(validation_dataloader)
validation_loss_values.append(val_loss)
print(f"Epoch {epoch+1}, Training loss: {train_loss:.4f}, Validation Loss: {val_loss:.4f}")

scheduler.step()
if val_loss < best_val_loss:
    best_val_loss = val_loss
    best_model = copy.deepcopy(model.state_dict()) # Save the best model
    counter_early_stop = 0 # Reset counter
    print("Validation loss decreased, saving model...")
else:
    counter_early_stop += 1
    print(f"Validation loss did not decrease, counter: {counter_early_stop}/{patience}")

if counter_early_stop >= patience or epoch == num_epochs - 1:

```

```
        early_stop = True
        break

torch.save(model.state_dict(), data_dir + '/model.pth')

plt.figure(figsize=(10, 5))
plt.plot(training_loss_values, label='Training Loss')
plt.plot(validation_loss_values, label='Validation Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.title('Loss vs. Epochs')
plt.legend()
plt.show()
```

Start Training...

EPOCH 1 of 200

Epoch 1, Training loss: 0.0099, Validation Loss: 0.0049  
Validation loss decreased, saving model...

EPOCH 2 of 200

Epoch 2, Training loss: 0.0017, Validation Loss: 0.0019  
Validation loss decreased, saving model...

EPOCH 3 of 200

Epoch 3, Training loss: 0.0012, Validation Loss: 0.0020  
Validation loss did not decrease, counter: 1/5

EPOCH 4 of 200

Epoch 4, Training loss: 0.0011, Validation Loss: 0.0021  
Validation loss did not decrease, counter: 2/5

EPOCH 5 of 200

Epoch 5, Training loss: 0.0009, Validation Loss: 0.0015  
Validation loss decreased, saving model...

EPOCH 6 of 200

Epoch 6, Training loss: 0.0009, Validation Loss: 0.0017  
Validation loss did not decrease, counter: 1/5

EPOCH 7 of 200

Epoch 7, Training loss: 0.0008, Validation Loss: 0.0023  
Validation loss did not decrease, counter: 2/5

EPOCH 8 of 200

Epoch 8, Training loss: 0.0008, Validation Loss: 0.0024  
Validation loss did not decrease, counter: 3/5

EPOCH 9 of 200

Epoch 9, Training loss: 0.0007, Validation Loss: 0.0014  
Validation loss decreased, saving model...

EPOCH 10 of 200

Epoch 10, Training loss: 0.0007, Validation Loss: 0.0027  
Validation loss did not decrease, counter: 1/5

EPOCH 11 of 200

Epoch 11, Training loss: 0.0006, Validation Loss: 0.0014  
Validation loss did not decrease, counter: 2/5

EPOCH 12 of 200

Epoch 12, Training loss: 0.0006, Validation Loss: 0.0013  
Validation loss decreased, saving model...

EPOCH 13 of 200

Epoch 13, Training loss: 0.0006, Validation Loss: 0.0015  
Validation loss did not decrease, counter: 1/5

EPOCH 14 of 200

Epoch 14, Training loss: 0.0006, Validation Loss: 0.0014  
Validation loss did not decrease, counter: 2/5

EPOCH 15 of 200

Epoch 15, Training loss: 0.0006, Validation Loss: 0.0015  
Validation loss did not decrease, counter: 3/5

EPOCH 16 of 200

Epoch 16, Training loss: 0.0005, Validation Loss: 0.0014  
Validation loss did not decrease, counter: 4/5

EPOCH 17 of 200

Epoch 17, Training loss: 0.0005, Validation Loss: 0.0012  
Validation loss decreased, saving model...

EPOCH 18 of 200

Epoch 18, Training loss: 0.0005, Validation Loss: 0.0014  
Validation loss did not decrease, counter: 1/5

EPOCH 19 of 200

Epoch 19, Training loss: 0.0005, Validation Loss: 0.0012  
Validation loss did not decrease, counter: 2/5

EPOCH 20 of 200

Epoch 20, Training loss: 0.0005, Validation Loss: 0.0015  
Validation loss did not decrease, counter: 3/5

EPOCH 21 of 200

Epoch 21, Training loss: 0.0005, Validation Loss: 0.0012  
Validation loss decreased, saving model...

EPOCH 22 of 200

Epoch 22, Training loss: 0.0005, Validation Loss: 0.0013  
Validation loss did not decrease, counter: 1/5

EPOCH 23 of 200

Epoch 23, Training loss: 0.0005, Validation Loss: 0.0012  
Validation loss decreased, saving model...

EPOCH 24 of 200

Epoch 24, Training loss: 0.0005, Validation Loss: 0.0012  
Validation loss did not decrease, counter: 1/5

EPOCH 25 of 200

Epoch 25, Training loss: 0.0005, Validation Loss: 0.0012  
Validation loss did not decrease, counter: 2/5

EPOCH 26 of 200

Epoch 26, Training loss: 0.0005, Validation Loss: 0.0012  
Validation loss decreased, saving model...

EPOCH 27 of 200

Epoch 27, Training loss: 0.0005, Validation Loss: 0.0013  
Validation loss did not decrease, counter: 1/5

EPOCH 28 of 200

Epoch 28, Training loss: 0.0005, Validation Loss: 0.0012  
Validation loss did not decrease, counter: 2/5

EPOCH 29 of 200


Epoch 29, Training loss: 0.0005, Validation Loss: 0.0012  
Validation loss did not decrease, counter: 3/5

EPOCH 30 of 200


Epoch 30, Training loss: 0.0005, Validation Loss: 0.0012  
Validation loss did not decrease, counter: 4/5

EPOCH 31 of 200

Epoch 31, Training loss: 0.0005, Validation Loss: 0.0012  
Validation loss did not decrease, counter: 5/5



```
In [ ]: plt.figure(figsize=(10, 5))
plt.plot(training_loss_values, label='Training Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.title('Loss vs. Epochs')
plt.legend()
plt.show()
```



```
In [ ]: train_batch_size = 10
        validation_batch_size = 10

train_dataset = ImageDataset(input_dir=colorization_data_dir, op="train", mask_json_path=mask_json, transform=transform)
validation_dataset = ImageDataset(input_dir=colorization_data_dir, op="val", mask_json_path=mask_json, transform=transform)
train_dataloader = DataLoader(train_dataset, batch_size=train_batch_size, shuffle=True)
validation_dataloader = DataLoader(validation_dataset, batch_size=validation_batch_size, shuffle=False)

device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
model = model.to(device)
outputs = model(next(iter(validation_dataloader))[0].to(device))


out_np = outputs[0].detach().cpu().numpy()
```



```
out_np = (out_np * 255).astype(np.uint8)

out_np = np.transpose(out_np, (1, 2, 0))

plt.imshow(out_np)
plt.axis('off')
plt.show()
```



```
In [ ]: ## Batch Size
train_batch_size = 10
validation_batch_size = 10

## Learning Rate
learning_rate = 0.001

# Epochs (Consider setting high and implementing early stopping)
num_epochs = 200
```

```

In [ ]: pretrain_model_path = torch.load(data_dir + '/model.pth')
n_classes = len(json.load(open(mask_json)))
model = UNET(n_classes)
device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
model = model.to(device)

model_dict = model.state_dict()

pretrain_dict = {k: v for k, v in pretrain_model_path.items() if k in model_dict and model_dict[k].size() :
model_dict.update(pretrain_dict)

model.load_state_dict(model_dict)

## Initialize Dataloaders
train_dataset=ImageDataset(input_dir=segmentation_data_dir, op="train", mask_json_path=mask_json, transform=
validation_dataset=ImageDataset(input_dir=segmentation_data_dir, op="val", mask_json_path=mask_json, transform=
test_dataset=ImageDataset(input_dir=segmentation_data_dir, op="test", mask_json_path=mask_json, transforms=
train_dataloader = DataLoader(train_dataset, batch_size=train_batch_size, shuffle=True)
validation_dataloader = DataLoader(validation_dataset, batch_size=validation_batch_size, shuffle=False)
test_dataloader = DataLoader(test_dataset, batch_size=1, shuffle=False)
## Initialize Optimizer and Learning Rate Scheduler
optimizer = torch.optim.Adam(model.parameters(),lr=learning_rate)
scheduler = torch.optim.lr_scheduler.StepLR(optimizer, step_size=10, gamma=0.1)

# implement early stopping
patience = 10
best_val_loss = float('inf')
counter_early_stop = 0
early_stop = False
training_loss_values = []
validation_loss_values = []
print("Start Training...")
for epoch in range(num_epochs):
    ##### Training #####
    print("\nEPOCH " +str(epoch+1)+" of "+str(num_epochs)+"\n")
    # TODO: Design your own training section
    model.train()
    train_loss = 0.0
    for images, masks in train_dataloader:

```

```

images = images.to(device)
masks = masks.to(device)

optimizer.zero_grad()
outputs = model(images)
# print(outputs.requires_grad)
loss = DICELoss(n_classes)(outputs, masks)
loss.backward()
optimizer.step()
train_loss += loss.item()

train_loss = train_loss / len(train_dataloader)
training_loss_values.append(train_loss)

##### Validation #####
# TODO: Design your own validation section
model.eval()
val_loss = 0.0
with torch.no_grad():
    for images, masks in validation_dataloader:
        images = images.to(device)
        masks = masks.to(device)
        outputs = model(images)
        # probs = torch.softmax(outputs, dim = 1)
        loss = DICELoss(n_classes).forward(outputs, masks)
        val_loss += loss.item()
val_loss = val_loss / len(validation_dataloader)
validation_loss_values.append(val_loss)
print(f"Epoch {epoch+1}, Training loss: {train_loss:.4f}, Validation Loss: {val_loss:.4f}")
if val_loss < best_val_loss:
    best_val_loss = val_loss
    best_model = copy.deepcopy(model.state_dict()) # Save the best model
    counter_early_stop = 0 # Reset counter
    print("Validation loss decreased, saving model...")
else:
    counter_early_stop += 1
    print(f"Validation loss did not decrease, counter: {counter_early_stop}/{patience}")

if counter_early_stop >= patience or epoch == num_epochs - 1:
    early_stop = True
    break

```

```
    scheduler.step()

plt.figure(figsize=(10, 5))
plt.plot(training_loss_values, label='Training Loss')
plt.plot(validation_loss_values, label='Validation Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.title('Loss vs. Epochs')
plt.legend()
plt.show()
```

Start Training...

EPOCH 1 of 200

Epoch 1, Training loss: 0.8570, Validation Loss: 0.7950  
Validation loss decreased, saving model...

EPOCH 2 of 200

Epoch 2, Training loss: 0.7780, Validation Loss: 0.7806  
Validation loss decreased, saving model...

EPOCH 3 of 200

Epoch 3, Training loss: 0.7267, Validation Loss: 0.7414  
Validation loss decreased, saving model...

EPOCH 4 of 200

Epoch 4, Training loss: 0.6772, Validation Loss: 0.7437  
Validation loss did not decrease, counter: 1/10

EPOCH 5 of 200

Epoch 5, Training loss: 0.6445, Validation Loss: 0.7653  
Validation loss did not decrease, counter: 2/10

EPOCH 6 of 200

Epoch 6, Training loss: 0.6176, Validation Loss: 0.7096  
Validation loss decreased, saving model...

EPOCH 7 of 200

Epoch 7, Training loss: 0.5733, Validation Loss: 0.6698  
Validation loss decreased, saving model...

EPOCH 8 of 200

Epoch 8, Training loss: 0.5500, Validation Loss: 0.6807  
Validation loss did not decrease, counter: 1/10

EPOCH 9 of 200

Epoch 9, Training loss: 0.5504, Validation Loss: 0.6984  
Validation loss did not decrease, counter: 2/10

EPOCH 10 of 200

Epoch 10, Training loss: 0.5344, Validation Loss: 0.6591  
Validation loss decreased, saving model...

EPOCH 11 of 200

Epoch 11, Training loss: 0.4741, Validation Loss: 0.6262  
Validation loss decreased, saving model...

EPOCH 12 of 200

Epoch 12, Training loss: 0.4698, Validation Loss: 0.6219  
Validation loss decreased, saving model...

EPOCH 13 of 200

Epoch 13, Training loss: 0.4476, Validation Loss: 0.6236  
Validation loss did not decrease, counter: 1/10

EPOCH 14 of 200

Epoch 14, Training loss: 0.4423, Validation Loss: 0.6176  
Validation loss decreased, saving model...

EPOCH 15 of 200

Epoch 15, Training loss: 0.4425, Validation Loss: 0.6187  
Validation loss did not decrease, counter: 1/10

EPOCH 16 of 200

Epoch 16, Training loss: 0.4437, Validation Loss: 0.6163  
Validation loss decreased, saving model...

EPOCH 17 of 200

Epoch 17, Training loss: 0.4390, Validation Loss: 0.6155  
Validation loss decreased, saving model...

EPOCH 18 of 200

Epoch 18, Training loss: 0.4315, Validation Loss: 0.6144  
Validation loss decreased, saving model...

EPOCH 19 of 200

Epoch 19, Training loss: 0.4345, Validation Loss: 0.6212  
Validation loss did not decrease, counter: 1/10

EPOCH 20 of 200

Epoch 20, Training loss: 0.4166, Validation Loss: 0.6120  
Validation loss decreased, saving model...

EPOCH 21 of 200

Epoch 21, Training loss: 0.4059, Validation Loss: 0.6103  
Validation loss decreased, saving model...

EPOCH 22 of 200

Epoch 22, Training loss: 0.4193, Validation Loss: 0.6097  
Validation loss decreased, saving model...

EPOCH 23 of 200

Epoch 23, Training loss: 0.3917, Validation Loss: 0.6097  
Validation loss decreased, saving model...

EPOCH 24 of 200

Epoch 24, Training loss: 0.4067, Validation Loss: 0.6102  
Validation loss did not decrease, counter: 1/10

EPOCH 25 of 200

Epoch 25, Training loss: 0.3973, Validation Loss: 0.6103  
Validation loss did not decrease, counter: 2/10

EPOCH 26 of 200

Epoch 26, Training loss: 0.4169, Validation Loss: 0.6083  
Validation loss decreased, saving model...

EPOCH 27 of 200

Epoch 27, Training loss: 0.4066, Validation Loss: 0.6091  
Validation loss did not decrease, counter: 1/10

EPOCH 28 of 200

Epoch 28, Training loss: 0.3942, Validation Loss: 0.6089  
Validation loss did not decrease, counter: 2/10

EPOCH 29 of 200

Epoch 29, Training loss: 0.3956, Validation Loss: 0.6100  
Validation loss did not decrease, counter: 3/10

EPOCH 30 of 200

Epoch 30, Training loss: 0.4009, Validation Loss: 0.6083  
Validation loss did not decrease, counter: 4/10

EPOCH 31 of 200

Epoch 31, Training loss: 0.4096, Validation Loss: 0.6085  
Validation loss did not decrease, counter: 5/10

EPOCH 32 of 200

Epoch 32, Training loss: 0.4004, Validation Loss: 0.6086  
Validation loss did not decrease, counter: 6/10

EPOCH 33 of 200

Epoch 33, Training loss: 0.4029, Validation Loss: 0.6088  
Validation loss did not decrease, counter: 7/10

EPOCH 34 of 200



Epoch 34, Training loss: 0.4018, Validation Loss: 0.6085  
Validation loss did not decrease, counter: 8/10

EPOCH 35 of 200

Epoch 35, Training loss: 0.3977, Validation Loss: 0.6082  
Validation loss decreased, saving model...

EPOCH 36 of 200

Epoch 36, Training loss: 0.3979, Validation Loss: 0.6090  
Validation loss did not decrease, counter: 1/10

EPOCH 37 of 200

Epoch 37, Training loss: 0.3952, Validation Loss: 0.6084  
Validation loss did not decrease, counter: 2/10

EPOCH 38 of 200

Epoch 38, Training loss: 0.3990, Validation Loss: 0.6079  
Validation loss decreased, saving model...

EPOCH 39 of 200

Epoch 39, Training loss: 0.4014, Validation Loss: 0.6085  
Validation loss did not decrease, counter: 1/10

EPOCH 40 of 200

Epoch 40, Training loss: 0.3929, Validation Loss: 0.6084  
Validation loss did not decrease, counter: 2/10

EPOCH 41 of 200

Epoch 41, Training loss: 0.3992, Validation Loss: 0.6071  
Validation loss decreased, saving model...

EPOCH 42 of 200

Epoch 42, Training loss: 0.4078, Validation Loss: 0.6085

Validation loss did not decrease, counter: 1/10

EPOCH 43 of 200

Epoch 43, Training loss: 0.3991, Validation Loss: 0.6086  
Validation loss did not decrease, counter: 2/10

EPOCH 44 of 200

Epoch 44, Training loss: 0.3819, Validation Loss: 0.6094  
Validation loss did not decrease, counter: 3/10

EPOCH 45 of 200

Epoch 45, Training loss: 0.3990, Validation Loss: 0.6090  
Validation loss did not decrease, counter: 4/10

EPOCH 46 of 200

Epoch 46, Training loss: 0.4043, Validation Loss: 0.6087  
Validation loss did not decrease, counter: 5/10

EPOCH 47 of 200

Epoch 47, Training loss: 0.3947, Validation Loss: 0.6089  
Validation loss did not decrease, counter: 6/10

EPOCH 48 of 200

Epoch 48, Training loss: 0.4003, Validation Loss: 0.6080  
Validation loss did not decrease, counter: 7/10

EPOCH 49 of 200

Epoch 49, Training loss: 0.3939, Validation Loss: 0.6078  
Validation loss did not decrease, counter: 8/10


EPOCH 50 of 200

Epoch 50, Training loss: 0.4160, Validation Loss: 0.6082  
Validation loss did not decrease, counter: 9/10

EPOCH 51 of 200

Epoch 51, Training loss: 0.3947, Validation Loss: 0.6074

Validation loss did not decrease, counter: 10/10



```
In [ ]: test_dice_score = dice_score_dataset(model, test_dataloader, n_classes, use_gpu=device.type == 'cuda')
        print(f"Epoch {epoch+1}, DICE score on the test set: {test_dice_score:.4f}")
```

Epoch 51, DICE score on the test set: 0.4896

## Problem 2: Transfer Learning

## Imports

```
In [ ]: ## Import VGG and FashionMNIST
import torch
from torchvision.models import vgg16
from torchvision.datasets import FashionMNIST
from torchvision import models, datasets, transforms
from torch.utils.data import DataLoader
from torch import nn
```

## Data Loading

```
In [ ]: ## Specify Batch Size
train_batch_size = 64
test_batch_size = 64

## Specify Image Transforms
img_transform = transforms.Compose([
    transforms.Resize((224,224)),
    transforms.Grayscale(num_output_channels= 3),
    transforms.ToTensor(),
    transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))
])

## Download Datasets
train_data = FashionMNIST('./data', transform=img_transform, download=True, train=True)
test_data = FashionMNIST('./data', transform=img_transform, download=True, train=False)

## Initialize Dataloaders
training_dataloader = DataLoader(train_data, batch_size=train_batch_size, shuffle=True)
test_dataloader = DataLoader(test_data, batch_size=test_batch_size, shuffle=True)
```

## Model Initialization and Training/Fine-tuning

Complete the rest of the assignment in the notebook below.

```
In [ ]: def init_random_vgg16():
    model = models.vgg16()
```

```

model.to('cuda')
return model

def finetune_vgg16():
    model = models.vgg16(pretrained=True)

    for param in model.features.parameters():
        param.requires_grad = False

    num_features = model.classifier[6].in_features
    # Freeze all but the last layer: randomly initialize the last layer of your network and fine-tune this
    model.classifier[6] = nn.Linear(num_features, len(train_data.classes))

    model.to('cuda')

    return model

```

In [ ]: *# TODO: Define your training loop, loss function, and optimizer and train your models*

```

def train_model(model, train_dataloader, criterion, optimizer, num_epochs=6):
    for epoch in range(num_epochs):
        model.train()
        running_loss = 0.0
        correct_predictions = 0

        for inputs, labels in train_dataloader:
            inputs, labels = inputs.to('cuda'), labels.to('cuda')
            optimizer.zero_grad()

            outputs = model(inputs)
            loss = criterion(outputs, labels)
            _, preds = torch.max(outputs, 1)

            loss.backward()
            optimizer.step()

            running_loss += loss.item() * inputs.size(0)
            correct_predictions += torch.sum(preds == labels.data)

        epoch_loss = running_loss / len(train_dataloader.dataset)
        epoch_acc = correct_predictions.double() / len(train_dataloader.dataset)

```

```
print(f'Epoch {epoch + 1}/{num_epochs}, Loss: {epoch_loss:.4f}, Accuracy: {epoch_acc:.4f}')

print('Training complete')
```

```
In [ ]: def test_model(model, test_dataloader, criterion, optimizer, num_epochs=6):
        model.eval()
        running_loss = 0.0
        correct_predictions = 0

        for inputs, labels in test_dataloader:
            inputs, labels = inputs.to('cuda'), labels.to('cuda')
            with torch.no_grad():
                outputs = model(inputs)
                loss = criterion(outputs, labels)
                _, preds = torch.max(outputs, 1)

            running_loss += loss.item() * inputs.size(0)
            correct_predictions += torch.sum(preds == labels.data)

        epoch_loss = running_loss / len(test_dataloader.dataset)
        epoch_acc = correct_predictions.double() / len(test_dataloader.dataset)

        print(f'Test Loss: {epoch_loss:.4f}, Test Acc: {epoch_acc:.4f}')
```

```
In [ ]: random_vgg16 = init_random_vgg16()
        finetuned_vgg16 = finetune_vgg16()
        criterion = nn.CrossEntropyLoss()
```

```
/usr/local/lib/python3.10/dist-packages/torchvision/models/_utils.py:208: UserWarning: The parameter 'pretrained' is deprecated since 0.13 and may be removed in the future, please use 'weights' instead.
  warnings.warn(
/usr/local/lib/python3.10/dist-packages/torchvision/models/_utils.py:223: UserWarning: Arguments other than a weight enum or `None` for 'weights' are deprecated since 0.13 and may be removed in the future. The current behavior is equivalent to passing `weights=VGG16_Weights.IMAGENET1K_V1`. You can also use `weights=VGG16_Weights.DEFAULT` to get the most up-to-date weights.
  warnings.warn(msg)
```

```
In [ ]: device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
        print(f"Current device: {device}")
```

Current device: cuda

```
In [ ]: optimizer_random = torch.optim.Adam(random_vgg16.parameters(), lr=0.001)
        train_model(random_vgg16, training_dataloader, criterion, optimizer_random, num_epochs=5)
```

```
Epoch 1/5, Loss: 1.2297, Accuracy: 0.6239
Epoch 2/5, Loss: 0.3174, Accuracy: 0.8840
Epoch 3/5, Loss: 0.2643, Accuracy: 0.9042
Epoch 4/5, Loss: 0.2397, Accuracy: 0.9121
Epoch 5/5, Loss: 0.2184, Accuracy: 0.9204
Training complete
```

```
In [ ]: test_model(random_vgg16, test_dataloader, criterion, optimizer_random, num_epochs=5)
```

```
Test Loss: 0.2340, Test Acc: 0.9156
```

```
In [ ]: print(random_vgg16(next(iter(test_dataloader))[0].to(device)).shape)
        print(next(iter(test_dataloader))[0].shape)
        print(next(iter(training_dataloader))[0].shape)
```

```
torch.Size([64, 1000])
torch.Size([64, 3, 224, 224])
torch.Size([64, 3, 224, 224])
```

```
In [ ]: optimizer_finetuned = torch.optim.Adam(finetuned_vgg16.classifier[6].parameters(), lr=0.001)
        train_model(finetuned_vgg16, training_dataloader, criterion, optimizer_finetuned, num_epochs=5)
```

```
Epoch 1/5, Loss: 0.5724, Accuracy: 0.7938
Epoch 2/5, Loss: 0.5141, Accuracy: 0.8149
Epoch 3/5, Loss: 0.5110, Accuracy: 0.8173
Epoch 4/5, Loss: 0.5183, Accuracy: 0.8191
Epoch 5/5, Loss: 0.5127, Accuracy: 0.8191
Training complete
```

```
In [ ]: test_model(finetuned_vgg16, test_dataloader, criterion, optimizer_finetuned, num_epochs=5)
```

```
Test Loss: 0.4118, Test Acc: 0.8498
```